

Machine Learning Project 1

Name - Aniket Nagdeo

```
In [1]: # Importing Packages
import numpy as np
import pandas as pd
import statistics
import matplotlib.pyplot as plt

# Importing seaborn sns
sns.set() # By default seaborn theme, scaling, and color palette
```

```
In [2]: df=pd.read_csv("C:/Users/Ankit/Desktop/bank-marketing.csv")
df.head()
```

	age	job	salary	marital	education	targeted	default	balance	housing	loan	contact	day	month	duration	campaign	pda
0	58	management	100000	married	tertiary	yes	no	2143	yes	no	unknown	5	may	261	1	
1	44	technician	60000	single	secondary	yes	no	29	yes	no	unknown	5	may	151	1	
2	33	entrepreneur	120000	married	secondary	yes	no	2	yes	yes	unknown	5	may	76	1	
3	47	blue-collar	20000	married	unknown	no	no	1506	yes	no	unknown	5	may	92	1	
4	33	unknown	0	single	unknown	no	no	1	no	no	unknown	5	may	198	1	

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4521 entries, 0 to 4520
Data columns (total 19 columns):
 # Column Non-Null Count Dtype
-----
 0 age 4521 non-null int64
 1 job 4521 non-null object
 2 salary 4521 non-null int64
 3 marital 4521 non-null object
 4 education 4521 non-null object
 5 targeted 4521 non-null object
 6 default 4521 non-null object
 7 balance 4521 non-null int64
 8 housing 4521 non-null object
 9 loan 4521 non-null object
 10 contact 4521 non-null object
 11 day 4521 non-null int64
 12 month 4521 non-null object
 13 duration 4521 non-null int64
 14 campaign 4521 non-null int64
 15 pdays 4521 non-null int64
 16 previous 4521 non-null int64
 17 poutcome 4521 non-null object
 18 response 4521 non-null object
dtypes: int64(8), object(11)
memory usage: 6.4+ MB
```

```
In [4]: df.describe()
```

	age	salary	balance	day	duration	campaign	pdays	previous
count	4521.000000	4521.000000	4521.000000	4521.000000	4521.000000	4521.000000	4521.000000	4521.000000
mean	40.936210	57066.17065	1362.272058	15.806419	258.63080	2.763841	40.197828	0.580323
std	10.618762	32085.71845	3044.765829	8.322476	257.527812	3.098021	100.128746	2.303441
min	18.000000	0.000000	-8019.000000	1.000000	0.000000	1.000000	-1.000000	0.000000
25%	33.000000	20000.000000	72.000000	8.000000	103.000000	1.000000	-1.000000	0.000000
50%	39.000000	60000.000000	448.000000	16.000000	180.000000	2.000000	-1.000000	0.000000
75%	48.000000	70000.000000	1428.000000	21.000000	319.000000	3.000000	-1.000000	0.000000
max	95.000000	120000.000000	102127.000000	31.000000	4918.000000	63.000000	871.000000	275.000000

```
In [5]: df.isnull().sum()
```

```
age 0
job 0
salary 0
marital 0
education 0
targeted 0
default 0
balance 0
housing 0
loan 0
contact 0
day 0
month 0
duration 0
campaign 0
pdays 0
previous 0
poutcome 0
response 0
dtype: int64
```

```
In [6]: df.shape()
```

```
Out[6]: (4521, 19)
```

```
In [7]: df.nunique()
```

```
age 77
job 12
salary 11
marital 3
education 4
targeted 2
default 2
balance 7168
housing 2
loan 2
contact 3
day 31
month 12
duration 1573
campaign 159
pdays 559
previous 41
poutcome 4
response 2
dtype: int64
```

Describe the pdays column, make note of the mean, median and minimum values. Anything fishy in the values?

```
In [8]: print('Mean is:',df['pdays'].mean())
print('Median is:',df['pdays'].median())
print('Mode is:',df['pdays'].mode()[0])
```

```
Mean is 40.19782796222158
Median is -1.0
Mode is -1
```

Describe the pdays column again, this time limiting yourself to the relevant values of pdays. How different are the mean and the median values?

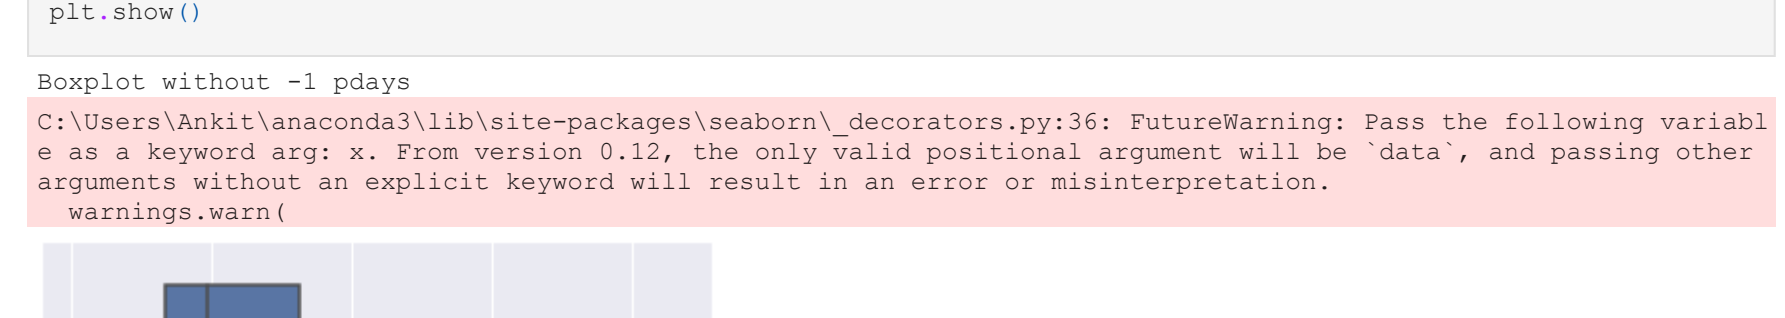
```
In [9]: # Since for relevant pdays, need to drop -1
df1 = df[df['pdays'] != -1]
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 8237 entries, 24060 to 43210
Data columns (total 19 columns):
 # Column Non-Null Count Dtype
-----
 0 age 8237 non-null int64
 1 job 8237 non-null object
 2 salary 8237 non-null int64
 3 marital 8237 non-null object
 4 education 8237 non-null object
 5 targeted 8237 non-null object
 6 default 8237 non-null object
 7 balance 8237 non-null int64
 8 housing 8237 non-null object
 9 loan 8237 non-null object
 10 day 8237 non-null int64
 11 month 8237 non-null object
 12 month 8237 non-null object
 13 duration 8237 non-null int64
 14 campaign 8237 non-null int64
 15 pdays 8237 non-null int64
 16 previous 8237 non-null int64
 17 poutcome 8237 non-null object
 18 response 8237 non-null object
dtypes: int64(8), object(11)
memory usage: 1.3+ MB
```

```
In [10]: print('Mean is:',df1['pdays'].mean())
print('Median is:',df1['pdays'].median())
print('Mode is:',df1['pdays'].mode()[0])
```

```
Mean is 224.5776916555496
Median is 194.0
Mode is 192
```

Plot a horizontal bar graph with the median values of balance for each education level value. Which group has the highest median?



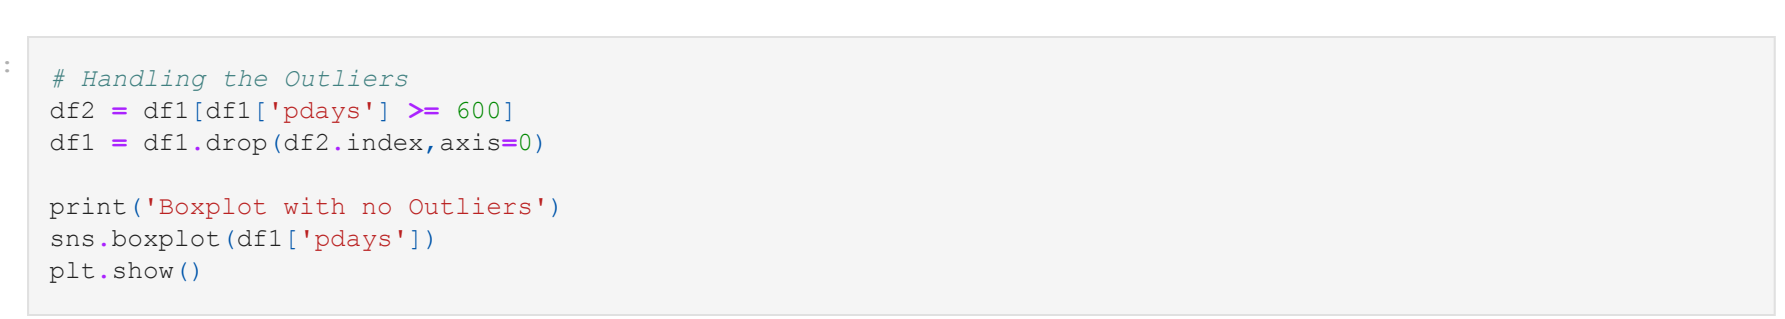
Unknown has highest median 782.0

Make a box plot for pdays. Do you see any outliers?



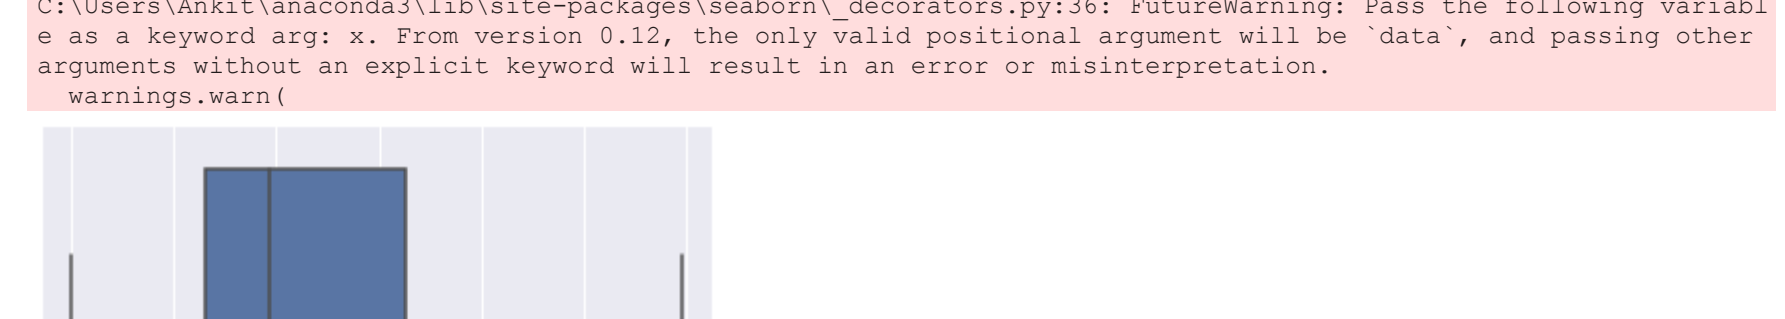
Boxplot without -1 pdays

```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



Boxplot with no Outliers

```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



The final goal is to make a predictive model to predict if the customer will respond positively to the campaign or not. The target variable is 'response'. First, perform bi-variate analysis to identify the features that are directly associated with the target variable. You can refer to the notebook we used for the EDA discussion.

EDA

Convert the response variable to a convenient form

```
In [15]: df1['response'] = np.where(df1['response'] == 'no', 0, 1)
```

```
In [16]: df1.head()
```

	age	job	salary	marital	education	targeted	default	balance	housing	loan	contact	day	month	duration	campaign	response
24060	32	admin.	50000	married	tertiary	yes	no	882	no	no	telephone	21	oct	39	1	0
24061	43	admin.	50000	single	secondary	yes	no	-247	yes	yes	telephone	21	oct	519	1	1
24064	33	services	70000	married	secondary	yes	no	3444	yes	no	telephone	21	oct	144	1	1
24072	36	management	100000	married	tertiary	yes	no	2415	yes	no	telephone	22	oct	73	1	1
24077	36	management	100000	married	tertiary	yes	no	0	yes	no	telephone	23	oct	140	1	1

Make suitable plots for associations with numerical features and categorical features'

```
In [17]: #Identifying Categorical and Numerical columns
cols = df1.columns.tolist()
num_cols = df1.get_numeric_data().columns.tolist()
cat_cols = list(set(cols) - set(num_cols))

print('Numerical Columns')
print(num_cols)
print('Categorical Columns')
print(cat_cols)
```

Numerical Columns

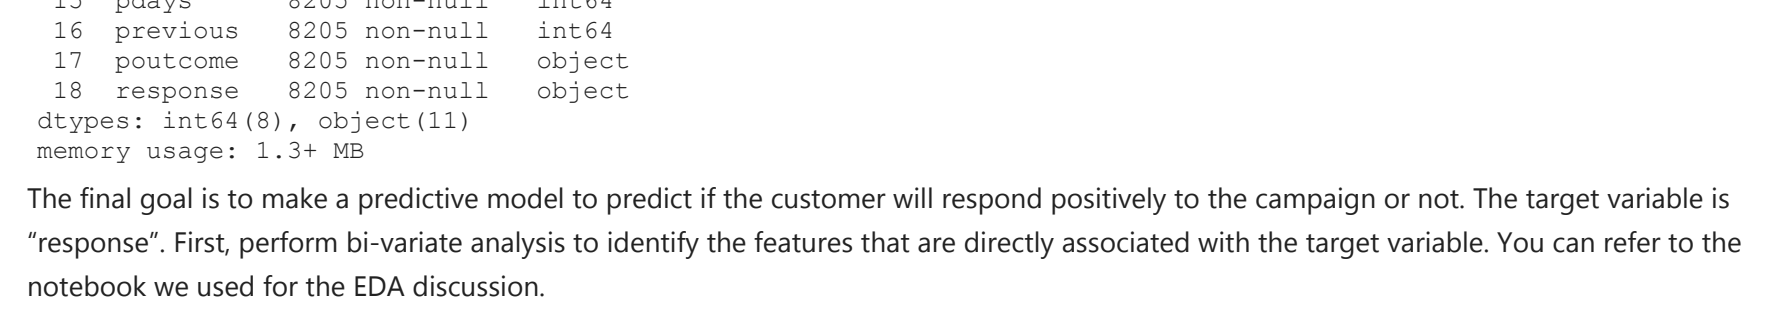
'age', 'salary', 'balance', 'day', 'duration', 'campaign', 'pdays', 'previous', 'response'

Categorical Columns

'default', 'job', 'marital', 'poutcome', 'month', 'targeted', 'loan', 'education', 'contact', 'housing'

```
In [18]: #Visualizing Numerical Variables with response
for i in df1(num_cols):
    sns.boxplot(df1['response'], df1[i])
    plt.show()
```

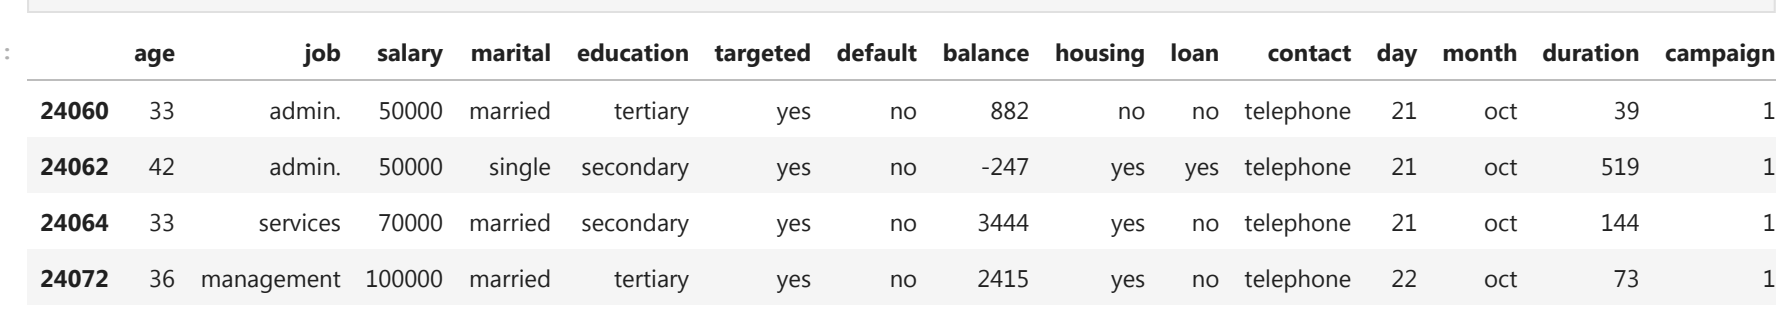
```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



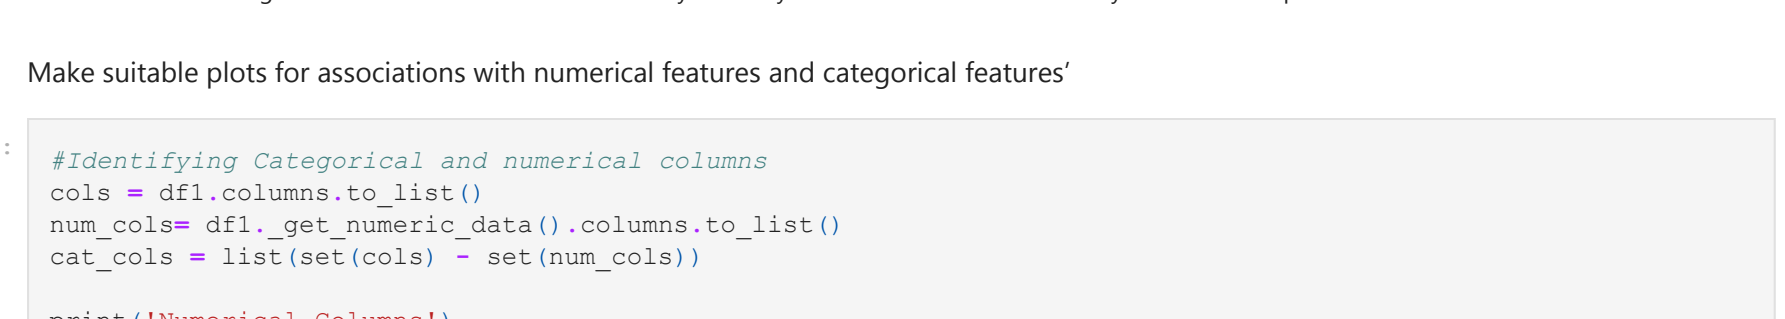
```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



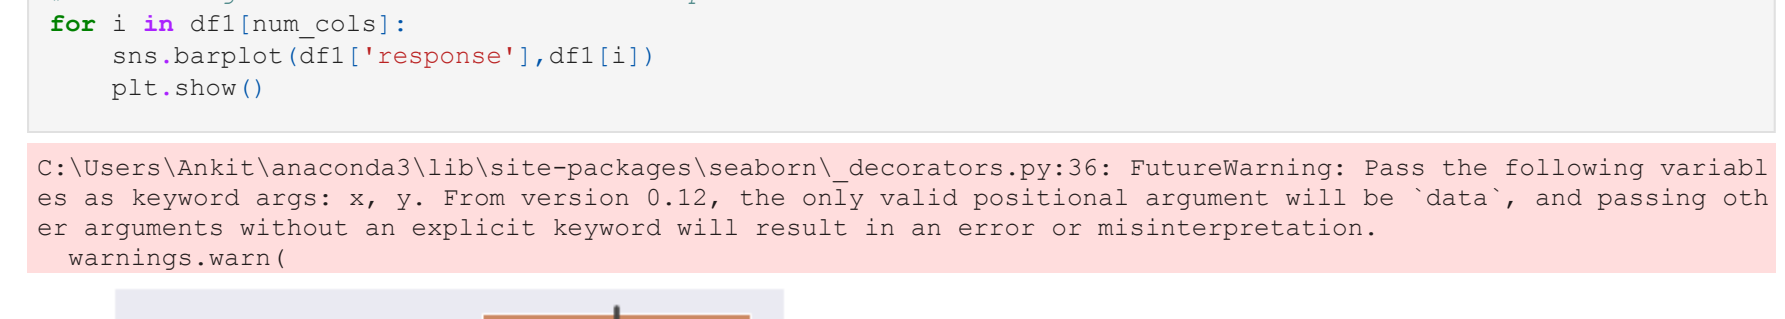
```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



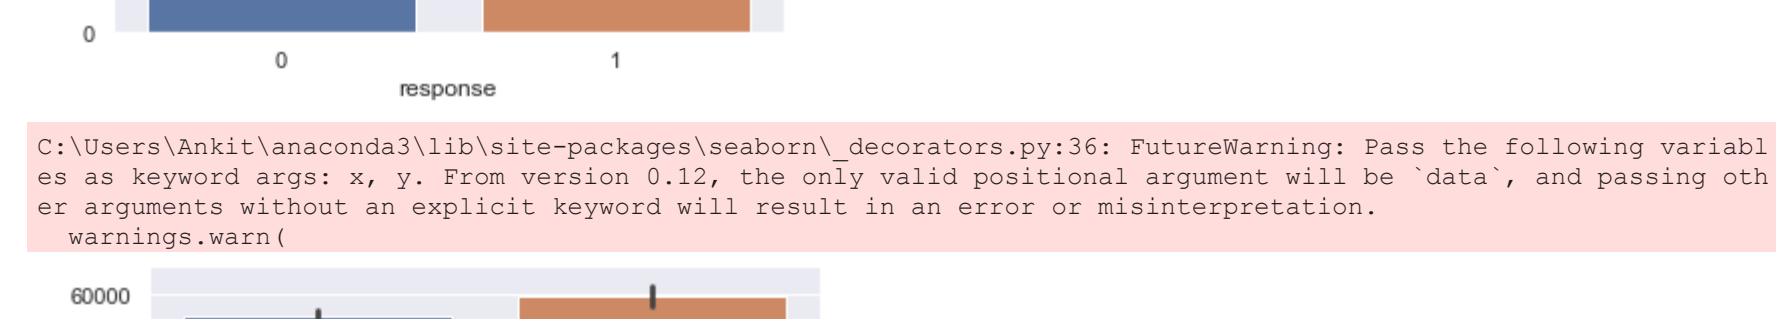
```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```




```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



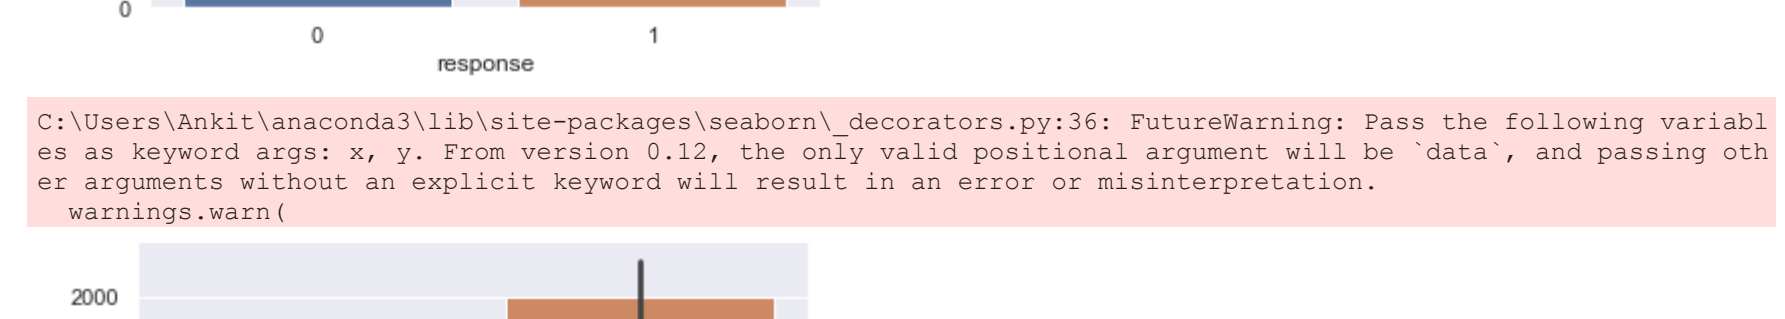
```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```




```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



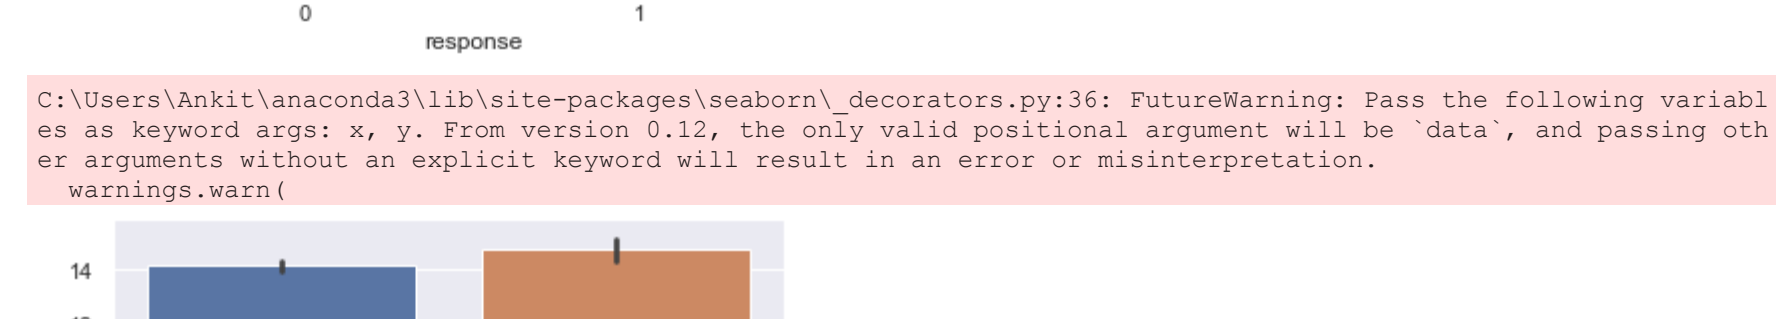
```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```




```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



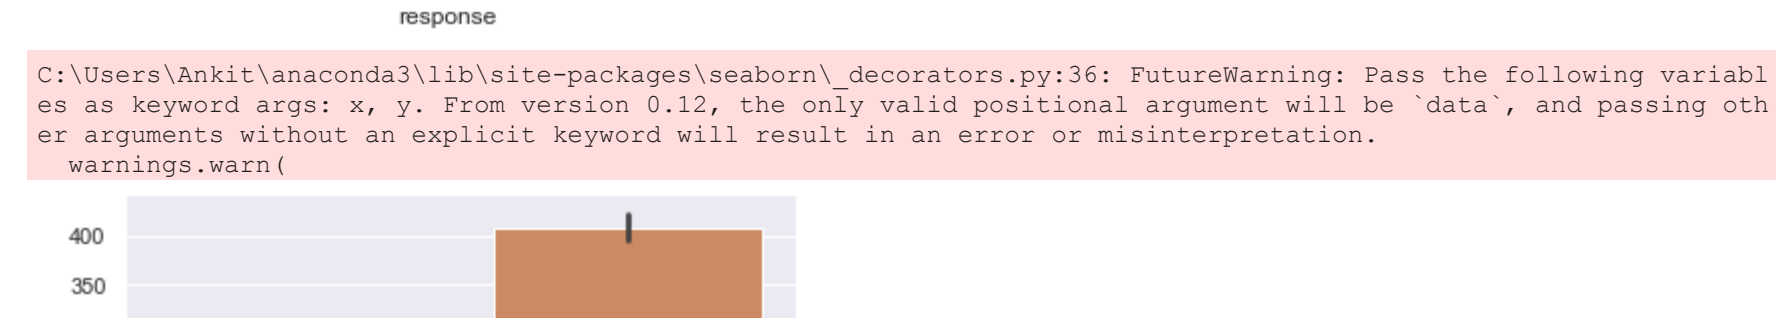
```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```




```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



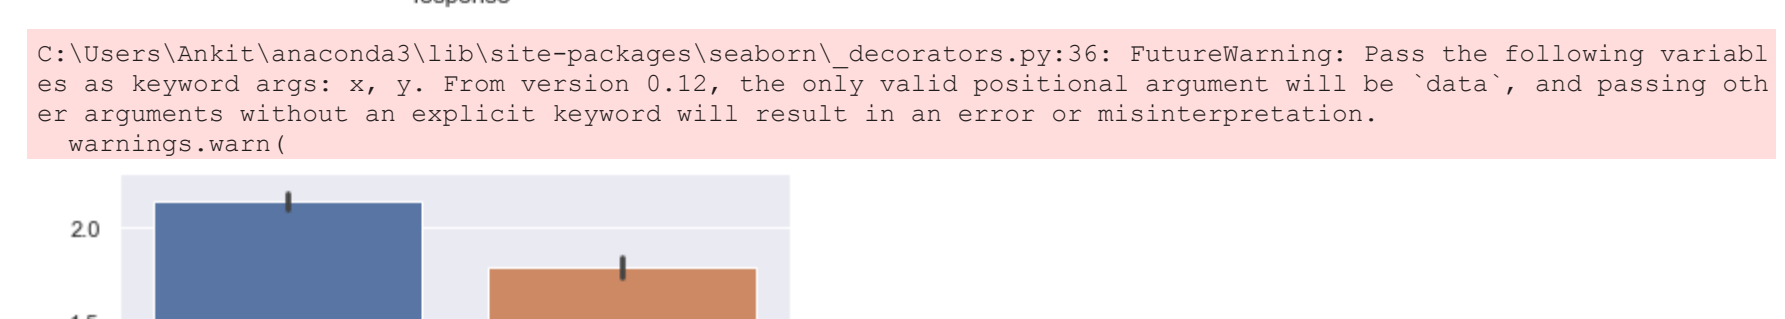
```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



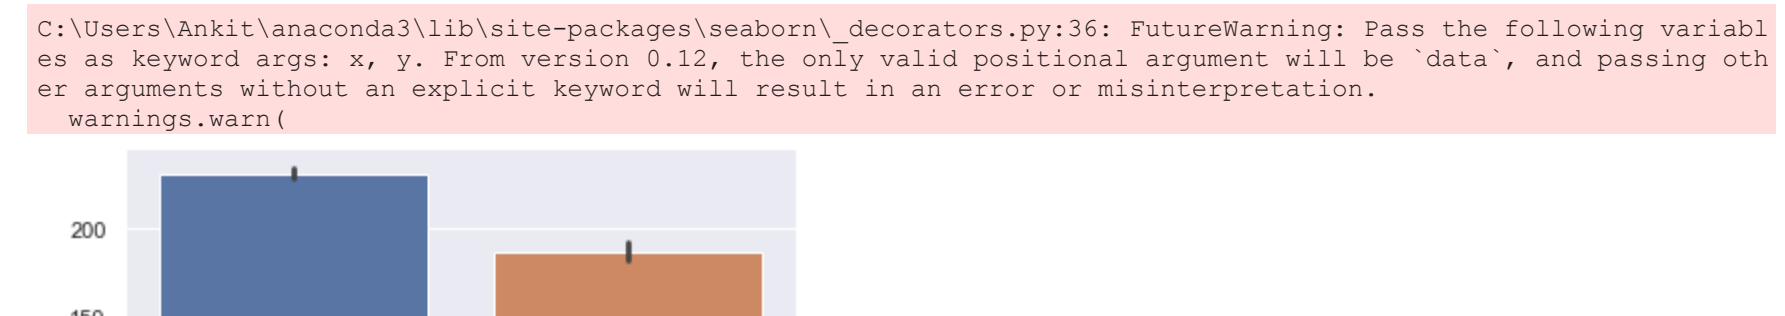
```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```




```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



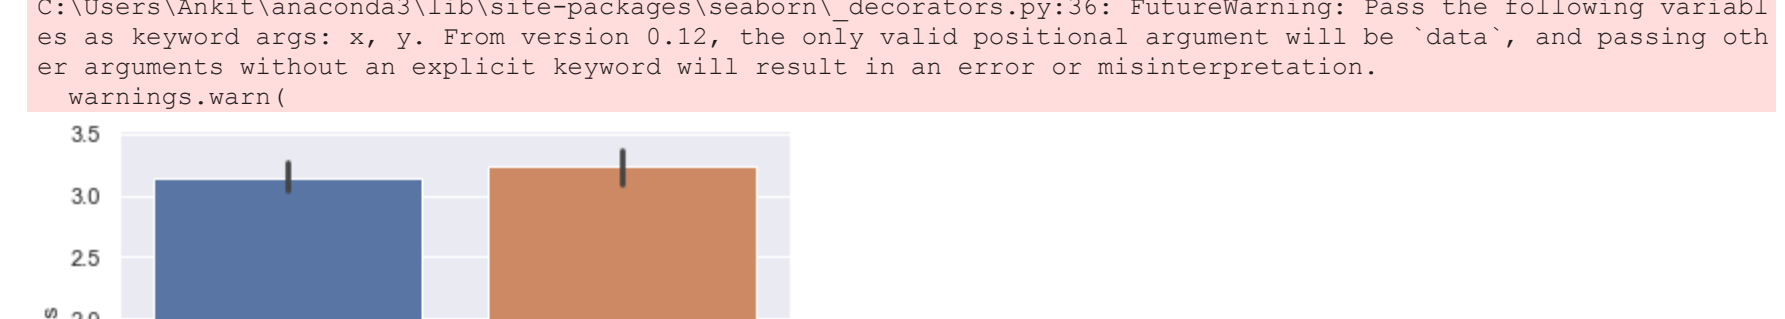
```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```




```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



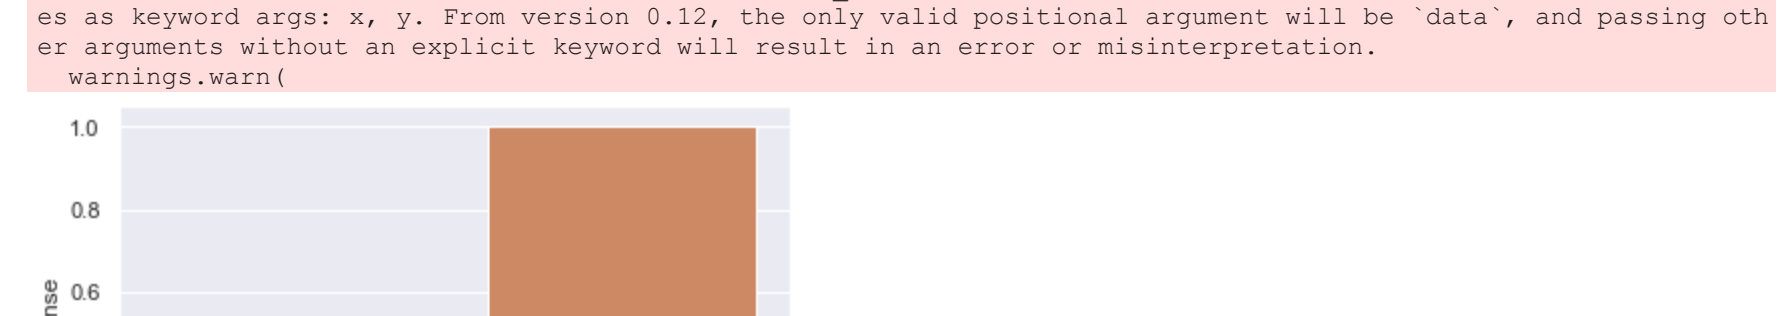
```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```




```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



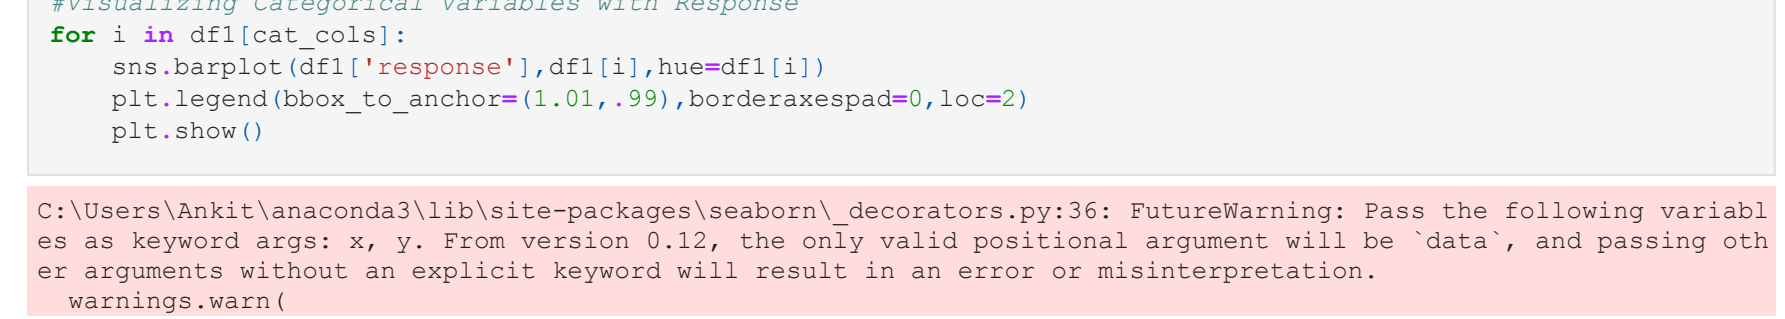
```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



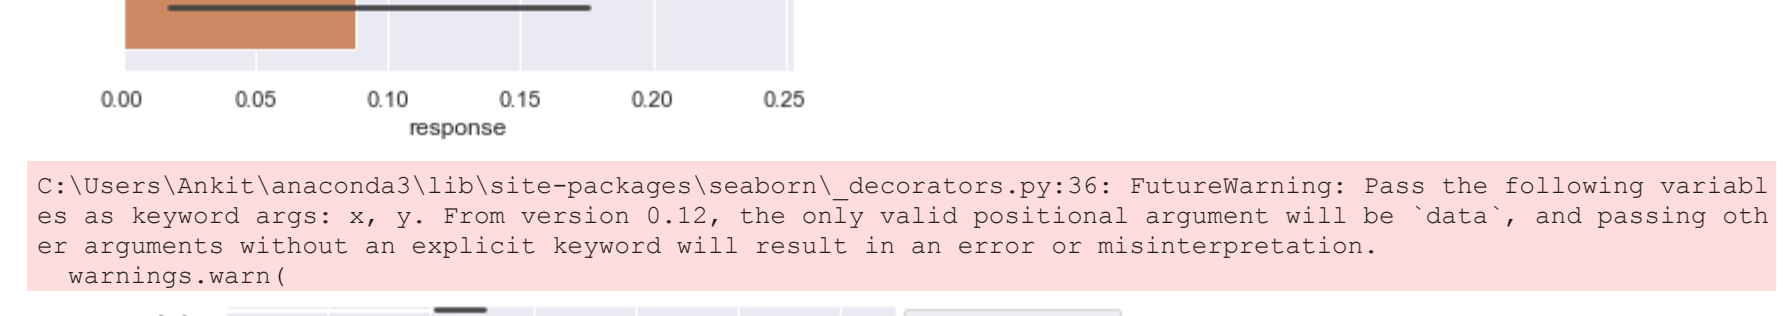
```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn()
```



```
C:\Users\Ankit\anaconda3\lib\site-packages\seaborn\decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be 'data', and passing other arguments without an explicit keyword will result in an error or
```



```
In [26]: # Actual Dummies
df1.drop(['poutcome', 'targeted', 'month', 'loan', 'education', 'housing', 'contact', 'marital', 'job', 'default', 'df_dummy = pd.concat([df1, df2], axis=1)
df_dummy.head()
```

	age	salary	balance	day	duration	campaign	pdays	previous	response	education_secondary	...	month_jul	month_jun	month_mar	r
0	33	50000	882	21	39	1	151	3	0	0	...	0	0	0	0
1	42	50000	-247	21	519	1	166	1	1	1	...	0	0	0	0
2	33	70000	3444	21	144	1	91	4	1	1	...	0	0	0	0
3	36	100000	2415	22	73	1	86	4	0	0	...	0	0	0	0
4	36	100000	0	23	140	1	143	3	1	1	...	0	0	0	0

5 rows x 32 columns

```
In [27]: ## Generally X - Independent variable, y - Dependent variable

X = df_dummy.drop(['response'], axis=1)
y = df_dummy['response']

print(X.head())
print(y.head())

age      salary  balance  day  duration  campaign  pdays  previous \
0      33    50000     882   21      39         1    151         3
1      42    50000    -247   21     519         1    166         1
2      33    70000    3444   21     144         1     91         4
3      36   100000    2415   22      73         1     86         4
4      36   100000         0   23     140         1    143         3

education_secondary  education_tertiary  ...  month_jul  month_jun \
0                0                1  ...         0         0
1                1                0  ...         0         0
2                1                0  ...         0         0
3                1                1  ...         0         0
4                0                1  ...         0         0

month_mar  month_may  month_nov  month_out  month_sep  poutcome_other \
0          0          0          0          1          0          0
1          0          0          0          1          0          1
2          0          0          0          1          0          0
3          0          0          0          1          0          1
4          0          0          0          1          0          0

poutcome_success  poutcome_unknown
0                0                0
1                0                0
2                0                0
3                0                0
4                0                0

[5 rows x 31 columns]
0      0
1      1
2      1
3      0
4      1
Name: response, dtype: int32
```

```
In [28]: # Train-Test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.3, random_state=1)
```

```
In [29]: # Checking data split
print(X_train.shape)
print(X_test.shape)
print(y_test.shape)

(5743, 31)
(5743,)
(2462, 31)
(2462,)
```

```
In [30]: # Scaling

# Applying Standard Scaling to make both of them comparable
from sklearn.preprocessing import StandardScaler
mm = StandardScaler()

X_train = mm.fit_transform(X_train)
X_test = mm.fit_transform(X_test)
```

```
In [31]: # Logistic Regression code
lg = LogisticRegression()
lg.fit(X_train, y_train)

y_train_pred = lg.predict(X_train)
y_test_pred = lg.predict(X_test)

print(y_train_pred)
print(y_test_pred)

[0 0 1 ... 1 0 0]
[0 0 0 ... 0 0 0]
```

```
In [32]: # Validating on Training
print(recall_score(y_train, y_train_pred))
print(precision_score(y_train, y_train_pred))
print(f1_score(y_train, y_train_pred))
print(accuracy_score(y_train, y_train_pred))
print("\n")

# Validating on Testing
print(recall_score(y_test, y_test_pred))
print(precision_score(y_test, y_test_pred))
print(f1_score(y_test, y_test_pred))
print(accuracy_score(y_test, y_test_pred))

0.5460725075528701
0.6999031945788964
0.6134912677131394
0.8413721051713132

0.5284697508896797
0.6971830985915493
0.6012145748987854
0.8399675060926076
```

```
In [33]: # Confusion Matrix for training data is:
print('Confusion Matrix for training data is:')
print(confusion_matrix(y_train, y_train_pred))
print("\n")
print('Confusion Matrix for testing data is:')
print(confusion_matrix(y_test, y_test_pred))

Confusion Matrix for training data is:
[[4109  310]
 [ 601  723]]

Confusion Matrix for testing data is:
[[1771  129]
 [ 265  297]]
```

Random Forest

```
In [37]: from sklearn.ensemble import RandomForestClassifier

rf_model = RandomForestClassifier(max_depth=14, min_samples_split=7)
rf_model.fit(X_train, y_train)

y_train_pred = rf_model.predict(X_train)
y_test_pred = rf_model.predict(X_test)

print(y_train_pred)
print(y_test_pred)

[0 0 1 ... 0 0 0]
[0 0 0 ... 0 0 0]
```

```
In [38]: # Validating on Training
print(recall_score(y_train, y_train_pred))
print(precision_score(y_train, y_train_pred))
print(f1_score(y_train, y_train_pred))
print(accuracy_score(y_train, y_train_pred))
print("\n")

# Validating on Testing
print(recall_score(y_test, y_test_pred))
print(precision_score(y_test, y_test_pred))
print(f1_score(y_test, y_test_pred))
print(accuracy_score(y_test, y_test_pred))

0.8225075528703906
0.9510917030567686
0.882138517628469
0.9493293618662024

0.5302491103202847
0.7376237623762376
0.646977256723778
0.8497156783103168
```

```
In [39]: # Confusion Matrix
print('Confusion Matrix for training data is:')
print(confusion_matrix(y_train, y_train_pred))
print("\n")
print('Confusion Matrix for testing data is:')
print(confusion_matrix(y_test, y_test_pred))

Confusion Matrix for training data is:
[[4363   56]
 [ 235 1089]]

Confusion Matrix for testing data is:
[[1794  106]
 [ 264  298]]
```